

## Рубежный контроль по дисциплине “Технологии машинного обучения”

Алексеев Андрей ИУ5-62Б

Вариант №1

### Задание

Для заданного набора данных (по варианту №1) постройте модели классификации или регрессии (в зависимости от конкретной задачи, рассматриваемой в наборе данных). Для построения моделей используйте методы 1 и 2 (Метод опорных векторов и Случайный лес). Оцените качество моделей на основе подходящих метрик качества (не менее двух метрик). Какие метрики качества Вы использовали и почему? Какие выводы Вы можете сделать о качестве построенных моделей? Для построения моделей необходимо выполнить требуемую предобработку данных: заполнение пропусков, кодирование категориальных признаков, и т.д.

```
{
  "authors": [
    {
      "name": "Алексеев Андрей Сергеевич"
    }
  ],
  "group": "ИУ5-62Б",
  "kernel_spec": {
    "name": "python3",
    "display_name": "Python 3 (ipykernel)",
    "language": "python"
  },
  "language_info": {
    "name": "python",
    "version": "3.9.7",
    "mimetype": "text/x-python",
    "codemirror_mode": {
      "name": "ipython",
      "version": 3
    },
    "pygments_lexer": "ipython3",
    "nbconvert_exporter": "python",
    "file_extension": ".py"
  },
  "title": "PKM2",
  "Вариант": "1"
}
```

```
{'authors': [{'name': 'Алексеев Андрей Сергеевич'}],
 'group': 'ИУ5-62Б',
 'kernel_spec': {'name': 'python3',
 'display_name': 'Python 3 (ipykernel)',
 'language': 'python'},
 'language_info': {'name': 'python',
 'version': '3.9.7',
 'mimetype': 'text/x-python',
 'codemirror_mode': {'name': 'ipython', 'version': 3},
 'pygments_lexer': 'ipython3',
 'nbconvert_exporter': 'python',
 'file_extension': '.py'},
 'title': 'PKM2',
 'Вариант': '1'}
```

```
import numpy as np
import pandas as pd
from typing import Dict, Tuple
from scipy import stats
from sklearn.datasets import load_iris, load_boston
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.neighbors import KNeighborsRegressor,
KNeighborsClassifier
```

```

from sklearn.metrics import accuracy_score, balanced_accuracy_score
from sklearn.metrics import plot_confusion_matrix
from sklearn.metrics import precision_score, recall_score, f1_score,
classification_report
from sklearn.metrics import confusion_matrix
from sklearn.preprocessing import StandardScaler, MinMaxScaler
from sklearn.metrics import mean_absolute_error, mean_squared_error,
mean_squared_log_error, median_absolute_error, r2_score
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import roc_curve, roc_auc_score
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
sns.set(style="ticks")

```

```

# https://scikit-learn.org/stable/datasets/index.html#iris-dataset

```

```

boston = load_boston()

```

```

# Наименования признаков

```

```

boston.feature_names

```

C:\Users\Админ\AppData\Local\Programs\Python\Python39\lib\site-packages\sklearn\utils\deprecation.py:87: FutureWarning: Function load\_boston is deprecated; `load\_boston` is deprecated in 1.0 and will be removed in 1.2.

The Boston housing prices dataset has an ethical problem. You can refer to the documentation of this function for further details.

The scikit-learn maintainers therefore strongly discourage the use of this dataset unless the purpose of the code is to study and educate about ethical issues in data science and machine learning.

In this special case, you can fetch the dataset from the original source::

```

import pandas as pd
import numpy as np

data_url = "http://lib.stat.cmu.edu/datasets/boston"
raw_df = pd.read_csv(data_url, sep="\s+", skiprows=22,
header=None)
data = np.hstack([raw_df.values[::2, :],
raw_df.values[1::2, :2]])
target = raw_df.values[1::2, 2]

```

Alternative datasets include the California housing dataset (i.e. :func:`~sklearn.datasets.fetch\_california\_housing`) and the Ames housing dataset. You can load the datasets as follows::

```
from sklearn.datasets import fetch_california_housing
housing = fetch_california_housing()
```

for the California housing dataset and::

```
from sklearn.datasets import fetch_openml
housing = fetch_openml(name="house_prices", as_frame=True)
```

for the Ames housing dataset.

```
warnings.warn(msg, category=FutureWarning)
```

```
array(['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS',
      'RAD',
      'TAX', 'PTRATIO', 'B', 'LSTAT'], dtype='<U7')
```

*# Значения признаков*

```
boston.data[:5]
```

```
array([[6.3200e-03, 1.8000e+01, 2.3100e+00, 0.0000e+00, 5.3800e-01,
        6.5750e+00, 6.5200e+01, 4.0900e+00, 1.0000e+00, 2.9600e+02,
        1.5300e+01, 3.9690e+02, 4.9800e+00],
       [2.7310e-02, 0.0000e+00, 7.0700e+00, 0.0000e+00, 4.6900e-01,
        6.4210e+00, 7.8900e+01, 4.9671e+00, 2.0000e+00, 2.4200e+02,
        1.7800e+01, 3.9690e+02, 9.1400e+00],
       [2.7290e-02, 0.0000e+00, 7.0700e+00, 0.0000e+00, 4.6900e-01,
        7.1850e+00, 6.1100e+01, 4.9671e+00, 2.0000e+00, 2.4200e+02,
        1.7800e+01, 3.9283e+02, 4.0300e+00],
       [3.2370e-02, 0.0000e+00, 2.1800e+00, 0.0000e+00, 4.5800e-01,
        6.9980e+00, 4.5800e+01, 6.0622e+00, 3.0000e+00, 2.2200e+02,
        1.8700e+01, 3.9463e+02, 2.9400e+00],
       [6.9050e-02, 0.0000e+00, 2.1800e+00, 0.0000e+00, 4.5800e-01,
        7.1470e+00, 5.4200e+01, 6.0622e+00, 3.0000e+00, 2.2200e+02,
        1.8700e+01, 3.9690e+02, 5.3300e+00]])
```

*# Значения целевого признака*

```
np.unique(boston.target)
```

```
array([ 5. ,  5.6,  6.3,  7. ,  7.2,  7.4,  7.5,  8.1,  8.3,  8.4,
        8.5,
        8.7,  8.8,  9.5,  9.6,  9.7, 10.2, 10.4, 10.5, 10.8, 10.9, 11.
        ,
        11.3, 11.5, 11.7, 11.8, 11.9, 12. , 12.1, 12.3, 12.5, 12.6,
       12.7,
        12.8, 13. , 13.1, 13.2, 13.3, 13.4, 13.5, 13.6, 13.8, 13.9, 14.
        ,
```

```

14.1, 14.2, 14.3, 14.4, 14.5, 14.6, 14.8, 14.9, 15. , 15.1,
15.2,
15.3, 15.4, 15.6, 15.7, 16. , 16.1, 16.2, 16.3, 16.4, 16.5,
16.6,
16.7, 16.8, 17. , 17.1, 17.2, 17.3, 17.4, 17.5, 17.6, 17.7,
17.8,
17.9, 18. , 18.1, 18.2, 18.3, 18.4, 18.5, 18.6, 18.7, 18.8,
18.9,
19. , 19.1, 19.2, 19.3, 19.4, 19.5, 19.6, 19.7, 19.8, 19.9, 20.
,
20.1, 20.2, 20.3, 20.4, 20.5, 20.6, 20.7, 20.8, 20.9, 21. ,
21.1,
21.2, 21.4, 21.5, 21.6, 21.7, 21.8, 21.9, 22. , 22.1, 22.2,
22.3,
22.4, 22.5, 22.6, 22.7, 22.8, 22.9, 23. , 23.1, 23.2, 23.3,
23.4,
23.5, 23.6, 23.7, 23.8, 23.9, 24. , 24.1, 24.2, 24.3, 24.4,
24.5,
24.6, 24.7, 24.8, 25. , 25.1, 25.2, 25.3, 26.2, 26.4, 26.5,
26.6,
26.7, 27. , 27.1, 27.5, 27.9, 28. , 28.1, 28.2, 28.4, 28.5,
28.6,
28.7, 29. , 29.1, 29.4, 29.6, 29.8, 29.9, 30.1, 30.3, 30.5,
30.7,
30.8, 31. , 31.1, 31.2, 31.5, 31.6, 31.7, 32. , 32.2, 32.4,
32.5,
32.7, 32.9, 33. , 33.1, 33.2, 33.3, 33.4, 33.8, 34.6, 34.7,
34.9,
35.1, 35.2, 35.4, 36. , 36.1, 36.2, 36.4, 36.5, 37. , 37.2,
37.3,
37.6, 37.9, 38.7, 39.8, 41.3, 41.7, 42.3, 42.8, 43.1, 43.5,
43.8,
44. , 44.8, 45.4, 46. , 46.7, 48.3, 48.5, 48.8, 50. ])
```

*# Размер выборки*

```
boston.data.shape, boston.target.shape
```

```
((506, 13), (506,))
```

*# Сформируем DataFrame*

```
data_url = "http://lib.stat.cmu.edu/datasets/boston"
```

```
raw_df = pd.read_csv(data_url, sep="\s+", skiprows=22, header=None)
```

```
data = np.hstack([raw_df.values[::2, :], raw_df.values[1::2, :2]])
```

```
target = raw_df.values[1::2, 2]
```

```
raw_df.rename(columns={0: 'CRIM'}, inplace=True)
```

```
raw_df.rename(columns={1: 'ZN'}, inplace=True)
```

```
raw_df.rename(columns={2: 'INDUS'}, inplace=True)
```

```
raw_df.rename(columns={3: 'CHAS'}, inplace=True)
```

```
raw_df.rename(columns={4: 'NOX'}, inplace=True)
```

```
raw_df.rename(columns={5: 'RM'}, inplace=True)
```

```

raw_df.rename(columns={6: 'AGE'}, inplace=True)
raw_df.rename(columns={7: 'DIS'}, inplace=True)
raw_df.rename(columns={8: 'RAD'}, inplace=True)
raw_df.rename(columns={9: 'TAX'}, inplace=True)
raw_df.rename(columns={10: 'PTRATIO'}, inplace=True)

```

*# Удаление строк, содержащих пустые значения*

```

raw_df_2 = raw_df.dropna(axis=0, how='any')
(raw_df.shape, raw_df_2.shape)

```

```

((1012, 11), (506, 11))

```

*# Проверим наличие пустых значений*

*# Цикл по колонкам датасета*

```

for col in raw_df_2.columns:

```

*# Количество пустых значений - все значения заполнены*

```

    temp_null_count = raw_df_2[raw_df_2[col].isnull()].shape[0]
    print('{} - {}'.format(col, temp_null_count))

```

```

CRIM - 0
ZN - 0
INDUS - 0
CHAS - 0
NOX - 0
RM - 0
AGE - 0
DIS - 0
RAD - 0
TAX - 0
PTRATIO - 0

```

```

raw_df_2['target'] = target
target1 = np.empty(len(raw_df_2['target']), dtype=np.int16)
j = 0
a = [0, 1, 2]
for i in raw_df_2['target']:
    if i <= 20 and i >= 5:
        target1[j] = 0
        j = j + 1
    if i <= 35 and i > 20:
        target1[j] = 1
        j = j + 1
    if i <= 50 and i > 35:
        target1[j] = 2
        j = j + 1
raw_df_2['target1'] = target1
raw_df_2.info()

```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 506 entries, 0 to 1010
Data columns (total 13 columns):

```

#	Column	Non-Null Count	Dtype
0	CRIM	506 non-null	float64
1	ZN	506 non-null	float64
2	INDUS	506 non-null	float64
3	CHAS	506 non-null	float64
4	NOX	506 non-null	float64
5	RM	506 non-null	float64
6	AGE	506 non-null	float64
7	DIS	506 non-null	float64
8	RAD	506 non-null	float64
9	TAX	506 non-null	float64
10	PTRATIO	506 non-null	float64
11	target	506 non-null	float64
12	target1	506 non-null	int16

dtypes: float64(12), int16(1)

memory usage: 52.4 KB

C:\Users\7272~1\AppData\Local\Temp\ipykernel\_2840\200341314.py:1:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation:

[https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
raw_df_2['target'] = target
```

C:\Users\7272~1\AppData\Local\Temp\ipykernel\_2840\200341314.py:15:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation:

[https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
raw_df_2['target1'] = target1
```

```
boston_X_train, boston_X_test, boston_y_train, boston_y_test =
train_test_split(
```

```
    raw_df_2.drop(['target1'], axis=1), raw_df_2['target1'],
    test_size=0.5, random_state=17)
```

```
def print_metrics(y_test, y_pred):
    rep = classification_report(y_test, y_pred, output_dict=True)
    print("weighted precision:", rep['weighted avg']['precision'])
    print("weighted recall:", rep['weighted avg']['recall'])
    print("weighted f1-score:", rep['weighted avg']['f1-score'])
    plt.figure(figsize=(4, 3))
    plt.title('Матрица ошибок')
    sns.heatmap(confusion_matrix(y_test, y_pred), annot=True,
cmap="Blues");
```

```

scaler = StandardScaler().fit(boston_X_train)
boston_X_train_scaled = pd.DataFrame(scaler.transform(boston_X_train),
columns=boston_X_train.columns)
boston_X_test_scaled = pd.DataFrame(scaler.transform(boston_X_test),
columns=boston_X_train.columns)
boston_X_train_scaled.describe()

```

	CRIM	ZN	INDUS	CHAS
NOX \				
count	2.530000e+02	2.530000e+02	2.530000e+02	253.000000
2.530000e+02				
mean	5.616939e-17	-7.021173e-18	-5.967997e-17	0.000000
2.281881e-16				
std	1.001982e+00	1.001982e+00	1.001982e+00	1.001982
1.001982e+00				
min	-4.896445e-01	-4.765149e-01	-1.585390e+00	-0.276759 -
1.411455e+00				
25%	-4.778570e-01	-4.765149e-01	-8.681584e-01	-0.276759 -
9.555808e-01				
50%	-4.474570e-01	-4.765149e-01	-2.232402e-01	-0.276759 -
1.600349e-01				
75%	1.132433e-01	1.022018e-01	1.017895e+00	-0.276759
6.712659e-01				
max	8.782363e+00	4.153219e+00	2.440552e+00	3.613247
2.816558e+00				

	RM	AGE	DIS	RAD
TAX \				
count	2.530000e+02	2.530000e+02	2.530000e+02	2.530000e+02
2.530000e+02				
mean	-1.011049e-15	-5.757362e-16	-2.106352e-16	-3.510587e-17
1.843058e-17				
std	1.001982e+00	1.001982e+00	1.001982e+00	1.001982e+00
1.001982e+00				
min	-3.742422e+00	-2.362153e+00	-1.258854e+00	-1.001348e+00 -
1.284044e+00				
25%	-5.981299e-01	-7.768940e-01	-8.206675e-01	-6.646074e-01 -
7.683696e-01				
50%	-1.456454e-01	2.894877e-01	-2.232697e-01	-5.523606e-01 -
4.670769e-01				
75%	5.432726e-01	9.014046e-01	6.066735e-01	1.580328e+00
1.485531e+00				
max	3.245951e+00	1.112534e+00	3.370167e+00	1.580328e+00
1.746265e+00				

	PTRATIO	target
count	2.530000e+02	2.530000e+02
mean	4.844610e-16	3.484257e-16
std	1.001982e+00	1.001982e+00
min	-2.667693e+00	-1.814881e+00



```

25%    -4.805161e-01 -6.609675e-01
50%     2.029766e-01 -1.561302e-01
75%     7.953370e-01  4.105239e-01
max      1.615528e+00  2.821379e+00

```

```

svm_model = SVC()
svm_model.fit(boston_X_train_scaled, boston_y_train)
boston_y_pred_svm = svm_model.predict(boston_X_test_scaled)
print_metrics(boston_y_test, boston_y_pred_svm)

```

```

weighted precision: 0.8705314009661835
weighted recall: 0.8656126482213439
weighted f1-score: 0.864150968697932

```



```

params = {'C': np.concatenate([np.arange(0.1, 2, 0.03), np.arange(2,
20, 1)])}
grid_cv = GridSearchCV(estimator=svm_model, param_grid=params, cv=10,
n_jobs=-1, scoring='f1_macro')
grid_cv.fit(boston_X_train_scaled, boston_y_train)
print(grid_cv.best_params_)

```

```

{'C': 11.0}

```

```

best_svm_model = grid_cv.best_estimator_
best_svm_model.fit(boston_X_train_scaled, boston_y_train)
boston_y_pred_svm = best_svm_model.predict(boston_X_test_scaled)
print_metrics(boston_y_test, boston_y_pred_svm)

```

```

weighted precision: 0.9129241554180029
weighted recall: 0.9130434782608695
weighted f1-score: 0.9127799736495389

```



```
rfc_model = RandomForestClassifier()
rfc_model.fit(boston_X_train, boston_y_train)
boston_y_pred_rfc = rfc_model.predict(boston_X_test)
print_metrics(boston_y_test, boston_y_pred_rfc)
```

weighted precision: 0.9884159318941929  
 weighted recall: 0.9881422924901185  
 weighted f1-score: 0.9877776823322747




```
params = {'n_estimators': [5, 10, 50, 100], 'max_features': [2, 3, 4],
          'criterion': ['gini', 'entropy'], 'min_samples_leaf': [1, 2, 3, 4, 5]}
grid_cv = GridSearchCV(estimator=rfc_model, param_grid=params, cv=10,
                        n_jobs=-1, scoring='f1_weighted')
grid_cv.fit(boston_X_train, boston_y_train)
print(grid_cv.best_params_)
```

```
{'criterion': 'gini', 'max_features': 4, 'min_samples_leaf': 1,
'n_estimators': 50}
```


```
best_rfc_model = grid_cv.best_estimator_  
best_rfc_model.fit(boston_X_train, boston_y_train)  
boston_y_pred_rfc = best_rfc_model.predict(boston_X_test)  
print_metrics(boston_y_test, boston_y_pred_rfc)
```

```
weighted precision: 0.9884159318941929  
weighted recall: 0.9881422924901185  
weighted f1-score: 0.9877776823322747
```



```
print_metrics(boston_y_test, boston_y_pred_svm)  
print_metrics(boston_y_test, boston_y_pred_rfc)
```

```
weighted precision: 0.9129241554180029  
weighted recall: 0.9130434782608695  
weighted f1-score: 0.9127799736495389  
weighted precision: 0.9884159318941929  
weighted recall: 0.9881422924901185  
weighted f1-score: 0.9877776823322747
```



## Вывод

Модели с подобранными гиперпараметрами оказались лучше базовых моделей. Обе конечные модели показали очень высокую точность прогноза, из-за того, что выбранный датасет является учебным. Из матриц ошибок видим, что модель метод опорных векторов совершила 22 неверных прогноза из 256, а модель случайный лес совершила 3 неверных прогноза из 256. Метрики показывают, что качества рассматриваемых моделей отличается.